


# Kafka Ecosystem



# Lesson Objectives

- Learn various tools and components around Kafka
- Confluent stack

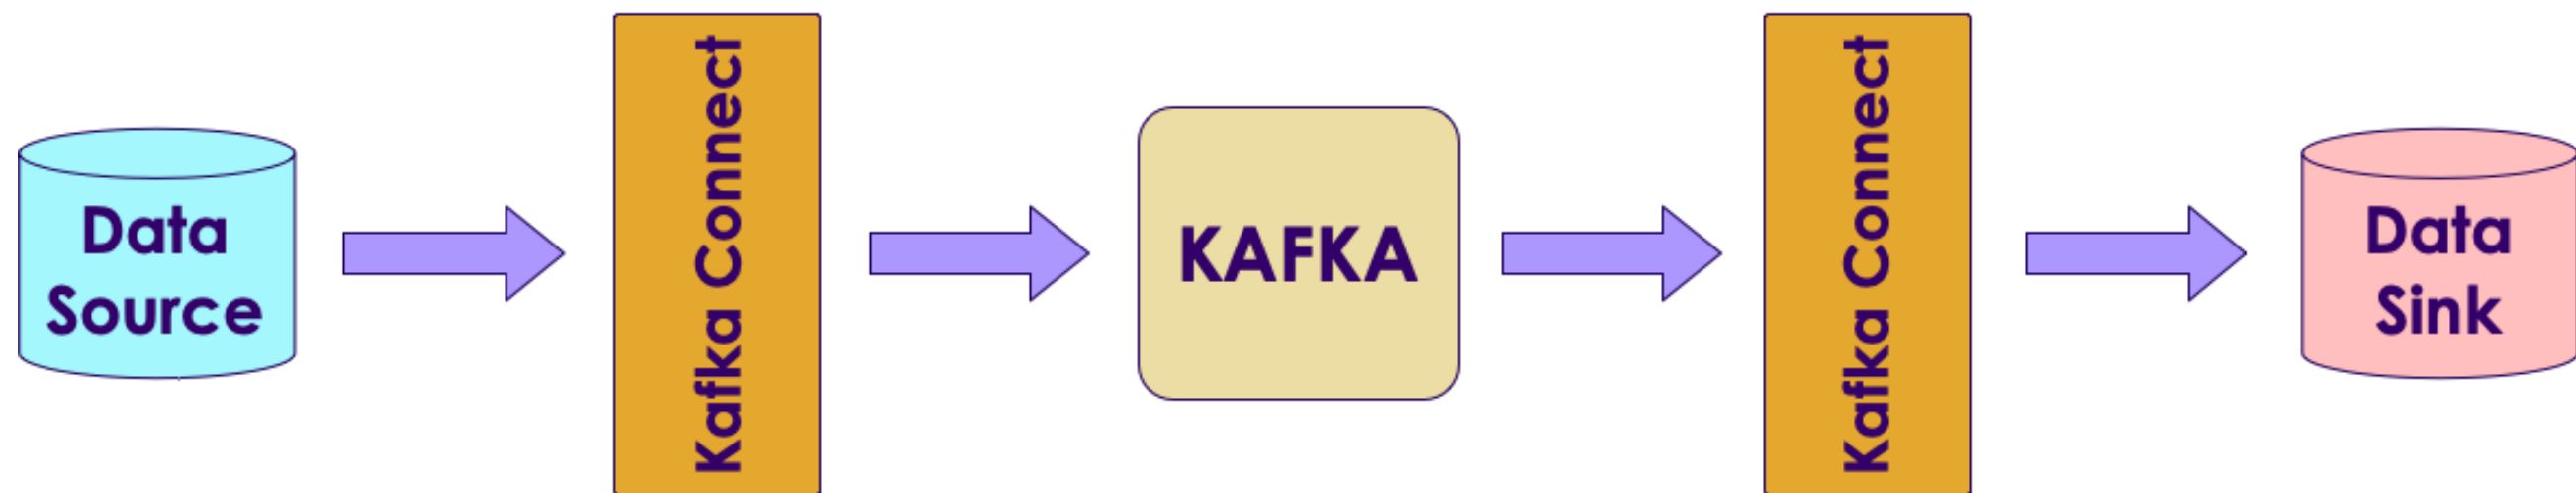
# Kafka Eco System

| Product          | Description   |
|------------------|---|
| Kafka Streams    | Build streaming applications easily                                     |
| Kafka Connect    | Move data between Kafka and other systems (DB / file system)            |
| Kafka Registry   | Metadata /schema store for data   |
| Kafka REST Proxy | REST interface into Kafka cluster.,Produce / Consume using RESTFUL APIs |
| Camus            | Kafka / HDFS integration  |

# Kafka Connect

# Kafka Connect

- Kafka Connect is a framework included in Apache Kafka that integrates Kafka with other systems.
- It's goal is to make it easy to add new systems to your scalable and secure stream data pipelines.



# Kafka Connectors (Supported by Confluent)

| Connector      | Description         | Supported by |
|----------------|---------------------|--------------|
| ActiveMQ       | Source only         | Confluent    |
| Amazon S3      | Sink                | Confluent    |
| Elastic Search | Sink                | Confluent    |
| HDFS           | Sink (Hadoop, Hive) | Confluent    |
| IBM MQ         | Source              | Confluent    |
| JDBC           | Source and Sink     | Confluent    |
| JMS            | Source              | Confluent    |

# Kafka Connectors (Supported by Vendors)

| Connector | Description     | Supported by |
|-----------|-----------------|--------------|
| Azure IoT | Source, IoT     | Microsoft    |
| Couchbase | Source and Sink | Couchbase    |
| SAP Hana  | Source and Sink | SAP          |
| Vertica   | Source and Sink | HP           |
| VoltDB    | Sink            | VoltDB       |

# Kafka Connectors (Supported by Community)

| Connector      | Description                           | Supported by |
|----------------|---------------------------------------|--------------|
| Amazon Kinesis | Sink (Amazon's managed queue service) | Community    |
| Apache Ignite  | Source and Sink (File System)         | Community    |
| Blockchain     | Source<br>Bitcoin,Blockchain          | Community    |
| Cassandra      | Sink<br>NoSQL                         | Community    |
| Github         | Source                                | Community    |
| Many more      |                                       |              |



# Kafka Connect Concepts

- **Connectors** - A logical process responsible for managing the copying of data between Kafka and another system.
- There are two types of connectors
  - **Source Connectors** import data from another system
  - **Sink Connectors** export data from Kafka
- **Tasks** - Unit of process that handles assigned set of work load by connectors. Connector configuration allows set to maximum number of tasks can be run by a connector.
- **Workers** - Unit of work that schedules connectors and tasks in a process.
- There are two main type of workers: **standalone** and **distributed**

# Standalone vs Distributed Workers

- Standalone Worker

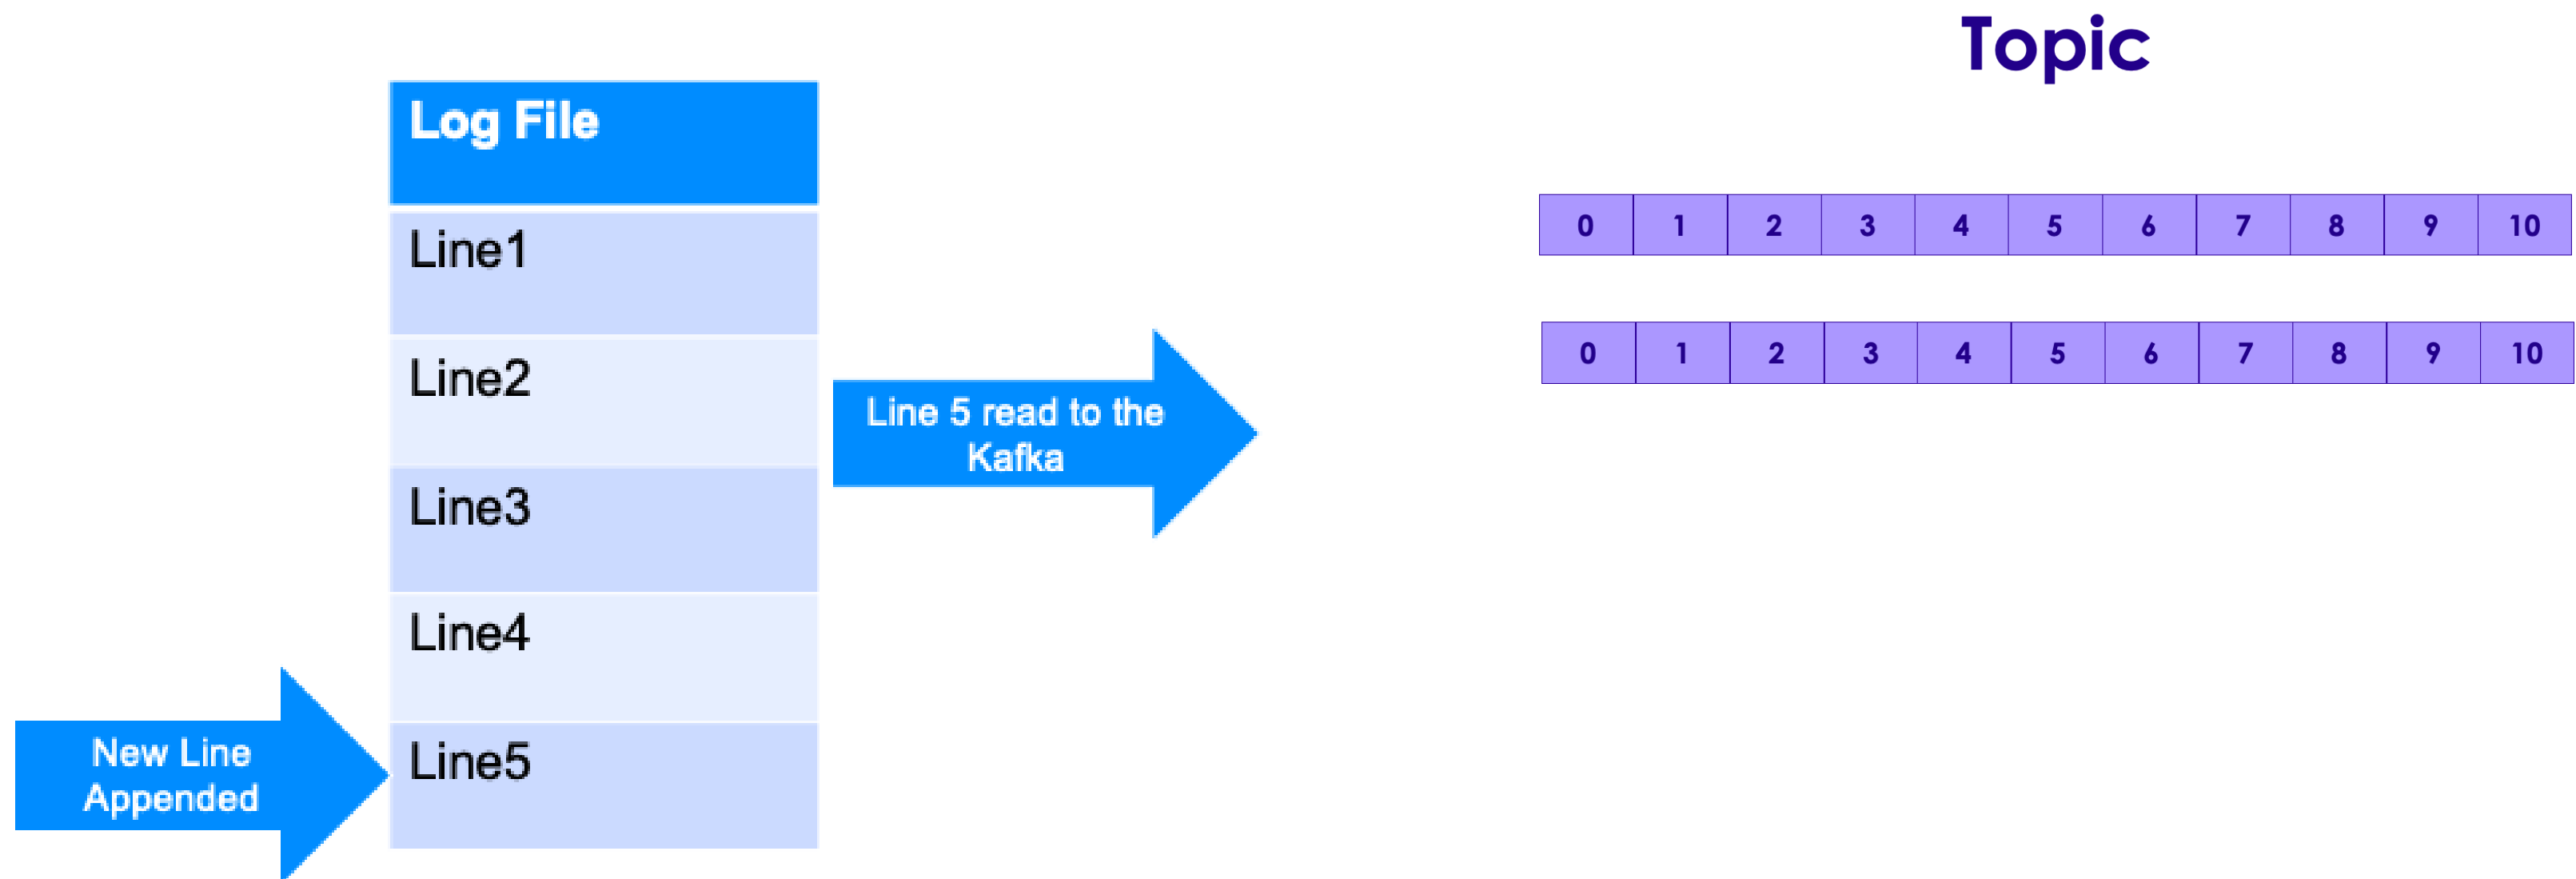
- Single process that executes all connectors and tasks
- Simple to configure
- Use for simple use cases or initial testing

- Distributed Worker

- Provides more scalability and fault tolerance
- Connectors and tasks are distributed between the workers automatically.

# File Connector

- Applications save logs into files on disk
- Streaming logfiles into kafka is a very common use case



# File Connector Configuration

- File source connector (File --> Kafka)

```
name=local-file-source
connector.class=FileStreamSource
tasks.max=1
file=/tmp/test.txt
topic=test-topic
```

- File sink connector (Kafka --> File)

```
name=local-file-sink
connector.class=FileStreamSink
tasks.max=1
file=/tmp/test.sink.txt
topics=test-topic
```

- Here is how to run it (in standalone mode)

```
bin/connect-standalone.sh \
  config/connect-standalone.properties \
  config/connect-file-source.properties
```

# HDFS Connector

- The HDFS connector exports data from Kafka into HDFS.
- Can also integrate with Hive, so data is readily available for querying using HiveQL
  - Hive tables are partitioned by Kafka topic
- Features
  - Exactly one delivery: Each Kafka message is only exported to HDFS once
  - Supports Avro and Parquet format
  - Secure data transport using Kerberos

# HDFS Connector

```
# hdfs-connector.properties  
  
name=hdfs-sink  
connector.class=io.confluent.connect.hdfs.HdfsSinkConnector  
tasks.max=1  
topics=test_hdfs_topic  
hdfs.url=hdfs://localhost:9000  
flush.size=3
```

- `flush.size` is how many records the connector need to write before invoking file commit
- Running it

```
$ confluent load hdfs-sink -d hdfs-connector.properties
```

- Checking data in HDFS

```
$ hadoop fs -ls /topics/test_hdfs/partition=0
```

```
/topics/test_hdfs/partition=0/test_hdfs+0+0000000000+0000000002.avro
```

# Running Connectors in Docker

- It is recommended to run Kafka Connect on containerized environments such as Kubernetes, Mesos, Docker Swarm, or YARN.
- Kafka Connect distributed mode exposes port 8083 by default to serve management REST interface.

```
$ docker run -d \  
  --name=kafka-connect \  
  --net=host \  
  -e CONNECT_BOOTSTRAP_SERVERS="kafka-broker:9092" \  
  -e CONNECT_GROUP_ID="group_1" \  
  -e CONNECT_CONFIG_STORAGE_TOPIC="kafka-connect-config" \  
  -e CONNECT_OFFSET_STORAGE_TOPIC="kafka-connect-offset" \  
  -e CONNECT_STATUS_STORAGE_TOPIC="kafka-connect-status" \  
  -e CONNECT_KEY_CONVERTER="org.apache.kafka.connect.json.JsonConverter" \  
  -e CONNECT_VALUE_CONVERTER="org.apache.kafka.connect.json.JsonConverter" \  
  -e CONNECT_INTERNAL_KEY_CONVERTER="org.apache.kafka.connect.json.JsonConverter" \  
  -e CONNECT_INTERNAL_VALUE_CONVERTER="org.apache.kafka.connect.json.JsonConverter" \  
  -e CONNECT_LOG4J_LOGGERS="io.debezium.connector.mysql=INFO" \  
  -v /opt/kafka-connect/jars:/etc/kafka-connect/jars \  
  --restart always \  
  confluentinc/cp-kafka-connect:3.3.0
```

# (Optional) Lab: Kafka Connect

- Overview:

- Use Kafka Connect to read data from a file

- Approximate Time:

- 20 - 30 mins

- Instructions:

- Please follow **kafka-connect** lab



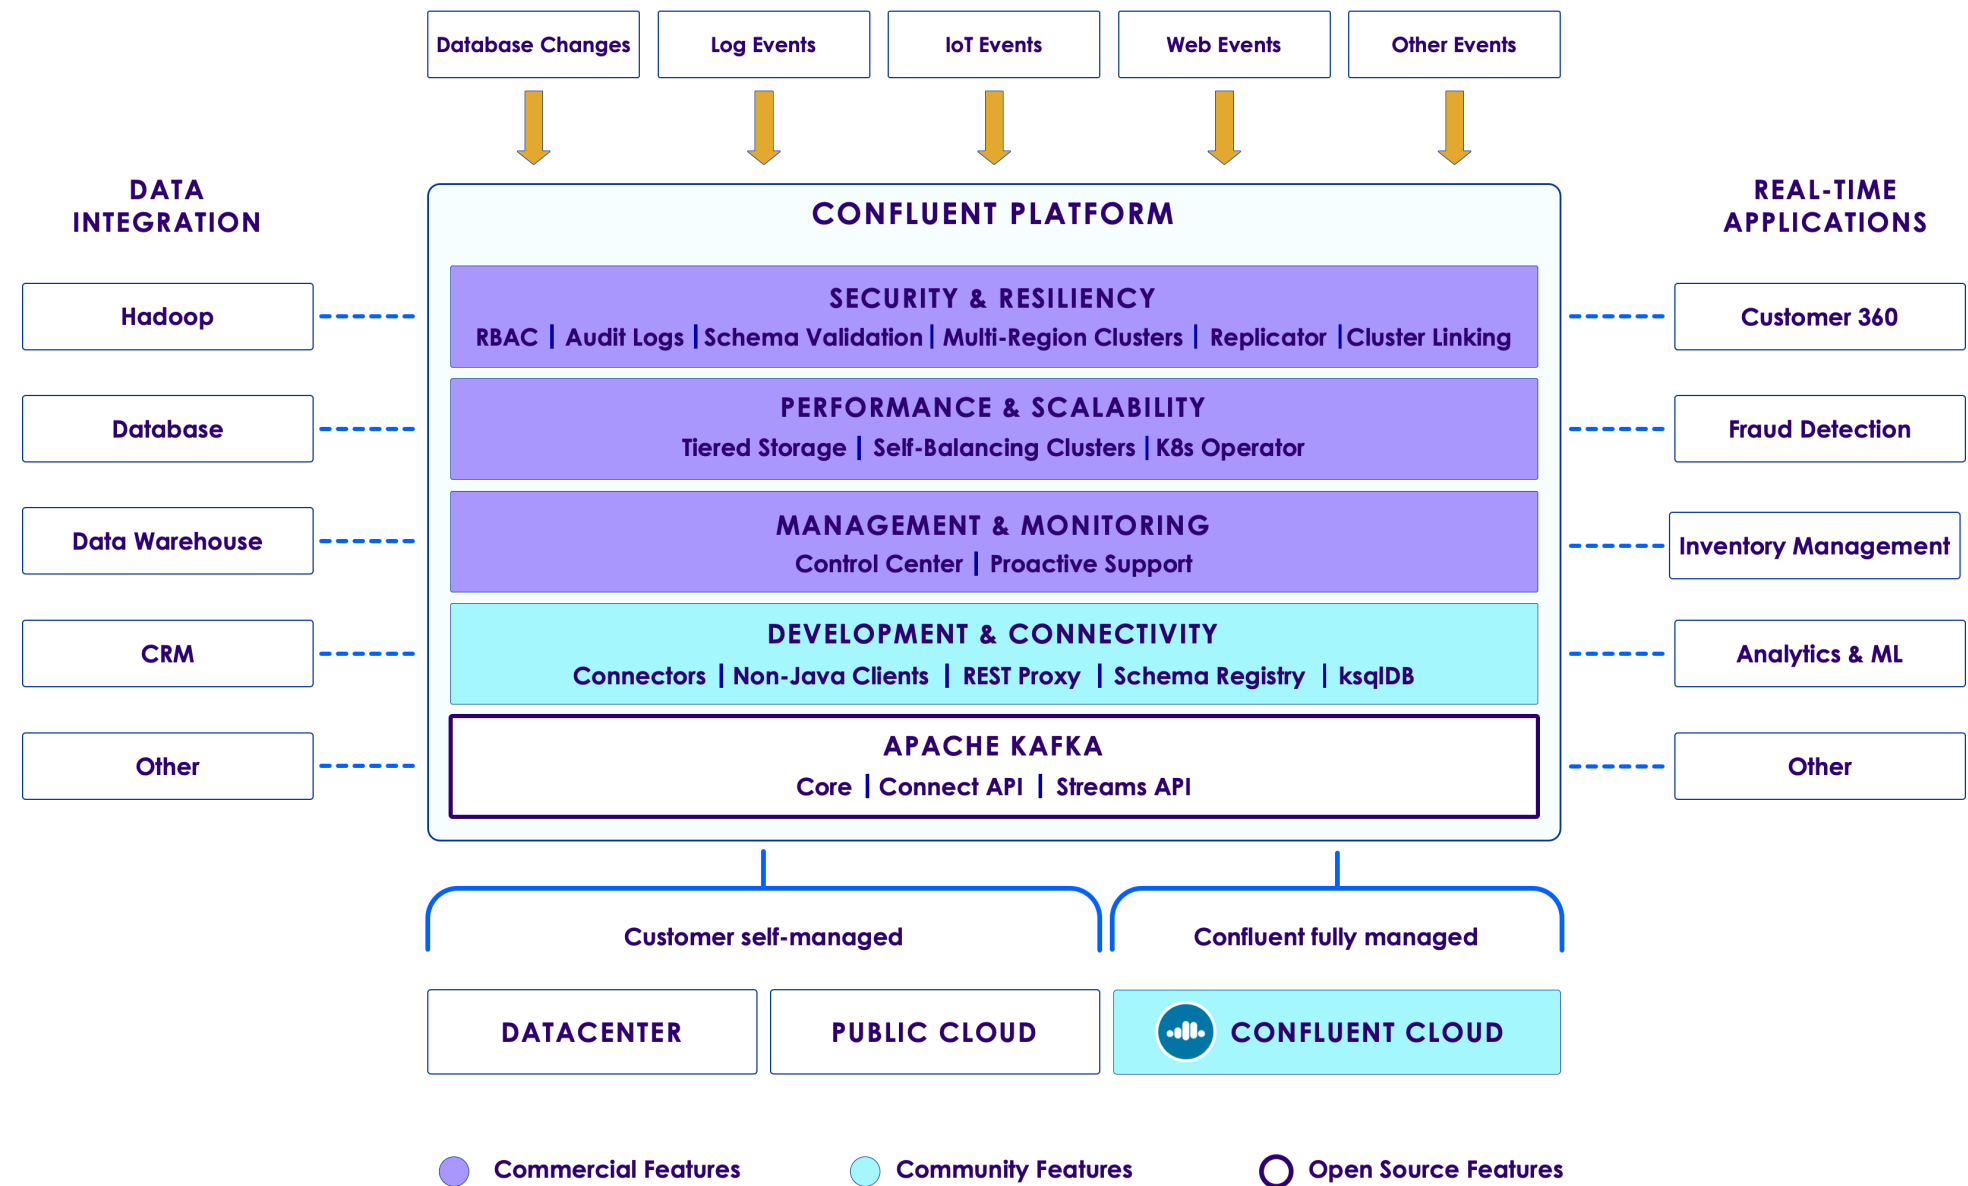


# Confluent Platform



# Confluent Platform

- Confluent platform has the following:
- Free features:
  - Apache Kafka
  - KSQL
  - Connectors (many types of databases)
  - Schema Registry
- Commercial features
  - Control Center
  - Replicator



# Lab: Setup Confluent Platform

- **Overview:**

- Download, install and start Confluent

- **Approximate Time:**

- 30 - 40 mins

- **Instructions:**

- CONFLUENT-1-SETUP



# Schema Registry

# Schemas Evolve in Real Life

- Let's say we have messages in the following format
- Version 1

| Id    | Type  | Success |
|-------|-------|---------|
| 12345 | Click | YES     |

- Version 2

| Id    | Type  | Success | Message        |
|-------|-------|---------|----------------|
| 12345 | Click | YES     | Page not found |

- Q: How will the consumer process this?

# Apache Avro

- Data serialization format
- Created for Hadoop project
- Language neutral; can be used from C, Java, Python ..etc
- Schema is described in JSON format
- Data is stored in binary format
- **Supports schema evolutions**
- Use KafkaAvroSerializer
- Apache Avro

# Avro Schema

## ■ Version 1

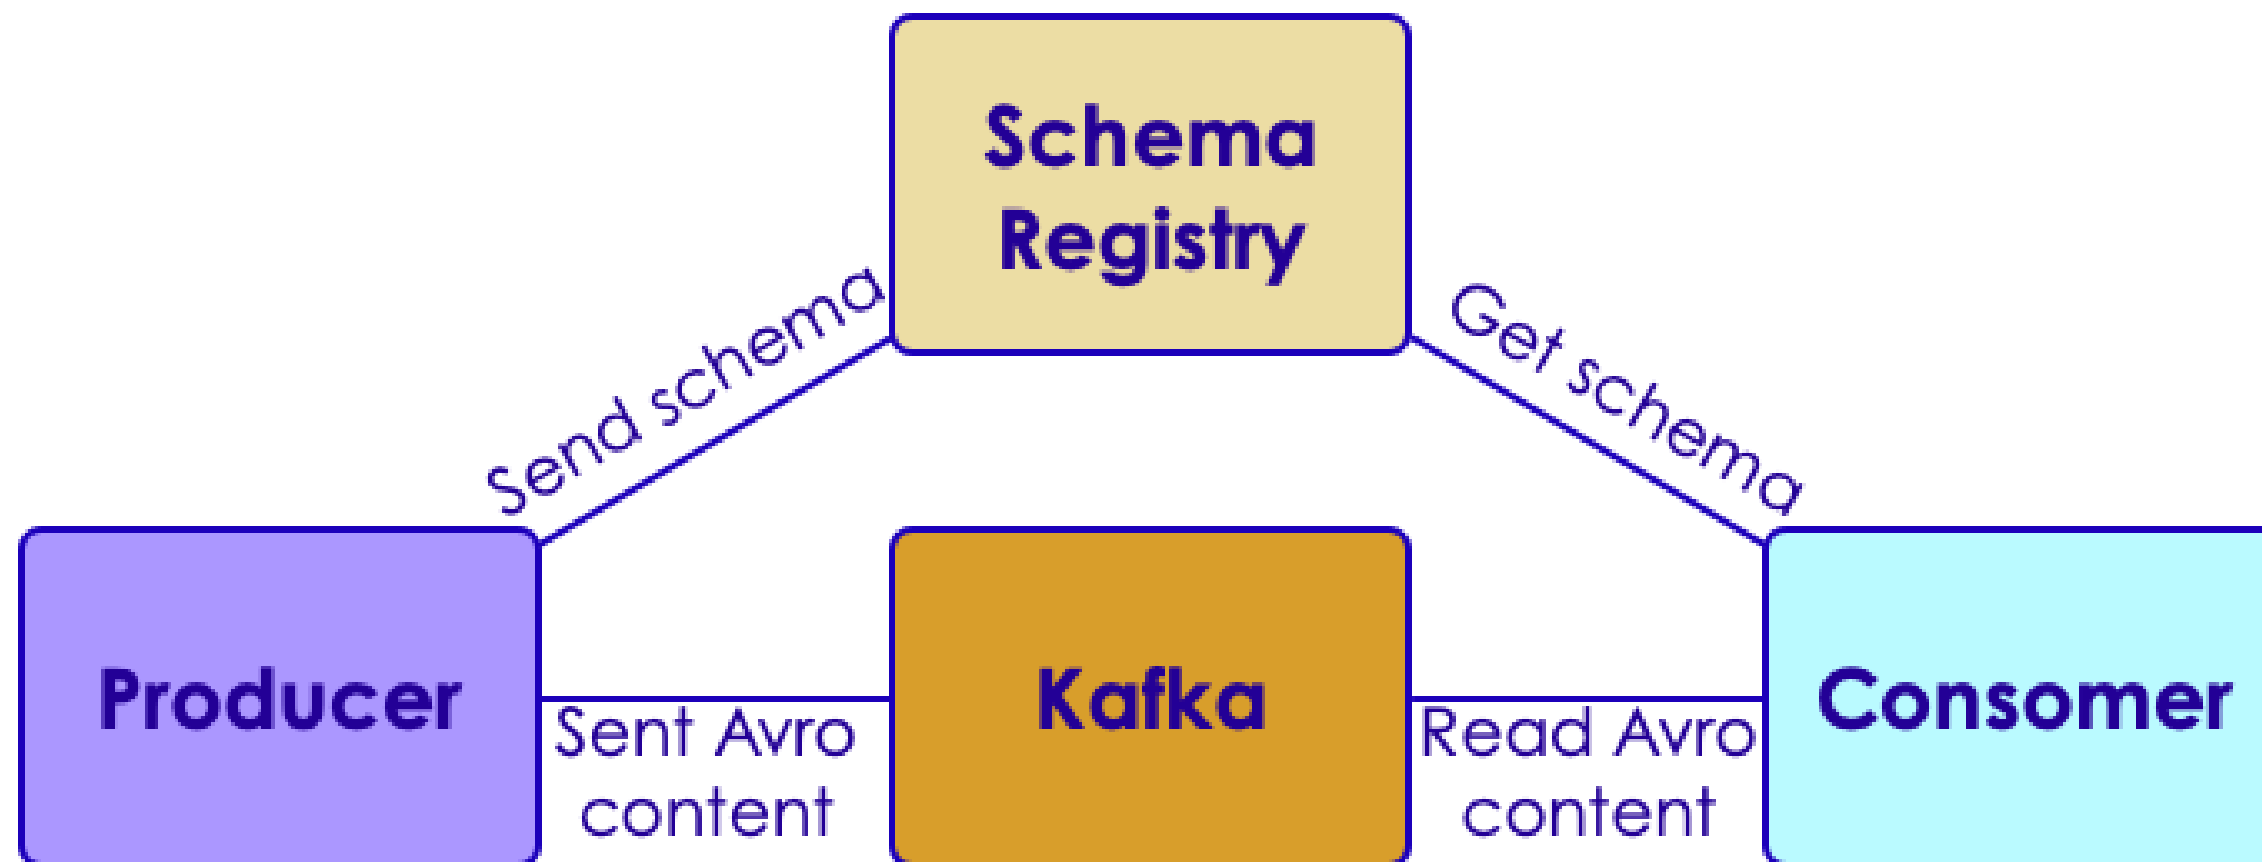
```
{
  "namespace": "com.example.videos",
  "type": "record",
  "name": "Event",
  "fields": [
    { "name": "id", "type": "int" },
    { "name": "type", "type": "string" },
    { "name": "success", "type": "string" }
  ]
}
```

## ■ Version 2

```
{
  "namespace": "com.example.videos",
  "type": "record",
  "name": "Event",
  "fields": [
    { "name": "id", "type": "int" },
    { "name": "type", "type": "string" },
    { "name": "success", "type": "string" },
    { "name": "message", "type": "string" } // <- new attribute
  ]
}
```

# Confluent Schema Registry

- Manages schemas and versions
- Provides REST API for interactions
- Works with Kafka seamlessly
- Open-source, downloadable as part of Confluent distribution
- [Documentation](#)





# Schema Registry Basics

- Schema

- Structure of an Avro data format

- Subject

- Scope of the schema. Subject name is derived from topic name by default
- Schemas can be registered under multiple subjects
  - Automated transparently from Producer as well

- Register a new schema:

```
curl -X POST -H "Content-Type: application/vnd.schemaregistry.v1+json" / \
--data '{"schema": "{/\"type/\": /\"string/\"}\"}' / \
http://localhost:8081/subjects/Kafka-value/versions {"id":1}
```

# Schema Registry Examples

- List all schemas under a subject

```
curl -X GET http://localhost:8081/subjects/Kafka-value/versions
```

- Fetch version 1 of the schema

```
curl -X GET http://localhost:8081/subjects/Kafka-value/versions/1  
# {"subject":"Kafka-value", "version":1, "id":1, "schema":"\"string\""}`
```

- Register the same schema under a different subject

```
curl -X POST -H "Content-Type: application/vnd.schemaregistry.v1+json" \  
  --data '{"schema": "{\"type\": \"string\"}'}' http://localhost:8081/subjects/Kafka2-value/versions  
# {"id":1}
```

# Lab: Use AVRO Schema Registry

- **Overview:**

- Setup and use Schema Registry

- **Approximate Time:**

- 30 - 40 mins

- **Instructions:**

- Avro lab 8.2



# Review and Q&A

- Let's go over what we have covered so far
- Any questions?

